

Contents

Overview.....	1
Stepup: vertex set, adjacency matrix, node features	1
Naive Approach: append node feaetures to adjacency matrix	1
3 problems: parameter size, graph size, node ordering	1
Adopt CNN	2
Problem	2
Solutions	2
Kipf & Welling, ICLR 2017	2
Explain – transformation & aggregation.....	2
Explain – each node has a computational graph	2
Many layers – k-hops.....	2
Neighbourhood aggregation.....	3
Transformation choice – average embedding	3
Model Parameters: W, B; neighborhood aggregation, transformation	3
Matrix Formulation - efficient sparse matrix	3
Averging of neighbour embedding	3
Re-writing update function in matrix form	3
How to train GNN: supervised & unsupervised setting.....	3
Unsupervised Training	4
Supervised Training	4
Cross entropy loss	4
Model design.....	4
Step 1,2 - neighbour aggregation & loss function	4
Step 3 – train on a batch.....	4
Step 4 - enerate embeddings for nodes as needed.....	4
Inductive capcability	4
New graph.....	4
New nodes.....	4
More.....	5
Permutation Invariance	5
Equivariant Property	5
GNNs subsume CNNs and Transformers	5
GNN vs. CNN	5

Transformer.....	5
Takeaway	5

Overview

Local network neighborhoods:

- Describe aggregation strategies
- Define computation graphs

Stacking multiple layers:

- Describe the model, parameters, training
- How to fit the model?
- Simple example for unsupervised and supervised training

■ Basics of neural networks

- Loss, Optimization, Gradient, SGD, non-linearity, MLP

■ Idea for Deep Learning for Graphs

- Multiple layers of embedding transformation
- At every layer, use the embedding at previous layer as the input
- Aggregation of neighbors and self-embeddings

■ Graph Convolutional Network

- Mean aggregation; can be expressed in matrix form

■ GNN is a general architecture

- CNN and Transformer can be viewed as a special GNN

Stepup: vertex set, adjacency matrix, node features

■ Assume we have a graph G :

- V is the **vertex set**
- A is the **adjacency matrix** (assume binary)
- $X \in \mathbb{R}^{m \times |V|}$ is a matrix of **node features**
- v : a node in V ; $N(v)$: the set of neighbors of v .
- **Node features:**
 - Social networks: User profile, User image
 - Biological networks: Gene expression profiles, gene functional information
 - When there is no node feature in the graph dataset:
 - Indicator vectors (one-hot encoding of a node)
 - Vector of constant 1: $[1, 1, \dots, 1]$


Naive Approach: append node feaetures to adjacency matrix

3 problems: parameter size, graph size, node ordering

1. Parameter size
 - a. One training example per node but for each node there are $N + X$ (node features) number of features
 - b. Training unstable / easy to overfit
2. Graph with different size
 - E.g. If 5 nodes, hard to fit in input size of 7
3. Node ordering
 - If the column order change, then the adjacency matrix changes
 - Rows & cols permuted thou the info is the same
 - For images, the ordering can be top left pixel to bottom right
 - but for graphs, there is no fix node ordering i.e. unclear how to sort the graph to put them as input in the matrix
 - Has to be **invariant to node ordering**

■ Join adjacency matrix and features

■ Feed them into a deep neural net:



■ Issues with this idea:

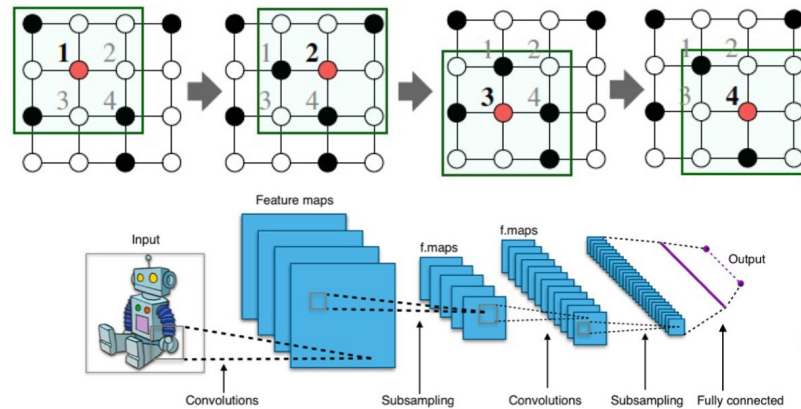
- $O(|V|)$ parameters
- Not applicable to graphs of different sizes
- Sensitive to node ordering

Adopt CNN

Goal. generalize convolutions beyond simple lattices & Leverage node features/attributes (e.g., text, images)

CNN = Sliding windows & locality.

CNN on an image:



Problem


- no fixed notion of locality or sliding window on the graph
 - L: covers 3 nodes
 - R: covers more nodes
- Graph is permutation invariant



Solutions

- Aggregate information about a node based on its neighbouring nodes

Single Convolutional neural network (CNN) layer with 3x3 filter:




Idea: transform information at the neighbors and combine it:

- Transform "messages" h_i from neighbors: $W_i h_i$
- Add them up: $\sum_i W_i h_i$

Kipf & Welling, ICLR 2017

- Neighbour nodes takes the message from the node and propagate. Steps
 - Determine node computation graph
 - Propagate & transform information

Idea: Node's neighborhood defines a computation graph




Learn how to propagate information across the graph to compute node features

Explain – transformation & aggregation

- To decide node A informatino, we collect from its neighbours {B, C, D}, in which their info are based on their neighbours.
- The message passing is then from the leaf to root
 - First, transform the info from leaf
 - Second, aggregate them in the parent node
 - Repeat

Intuition: Nodes aggregate information from their neighbors using neural networks

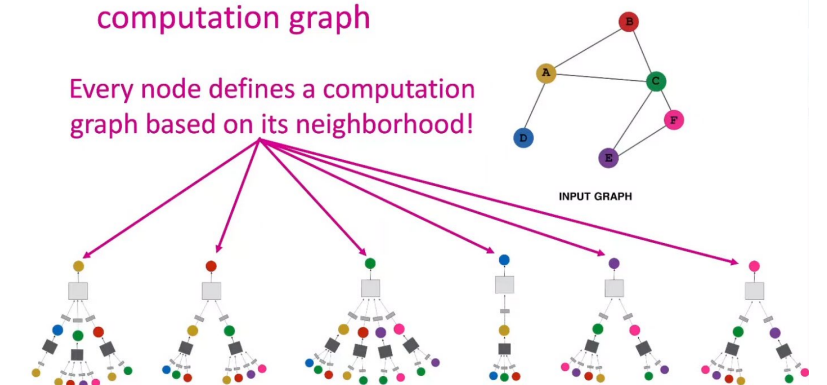


Explain – each node has a computational graph

- Every node has its own computation graph / architecture
- The structure depends on other structure
- Different to classical DL

Intuition: Network neighborhood defines a computation graph


Every node defines a computation graph based on its neighborhood!



Many layers – k-hops

Model can be of arbitrary depth:


- Nodes have embeddings at each layer
- Layer-0 embedding of node u is its input feature, x_u
- Layer- k embedding gets information from nodes that are K hops away



Neighbourhood aggregation

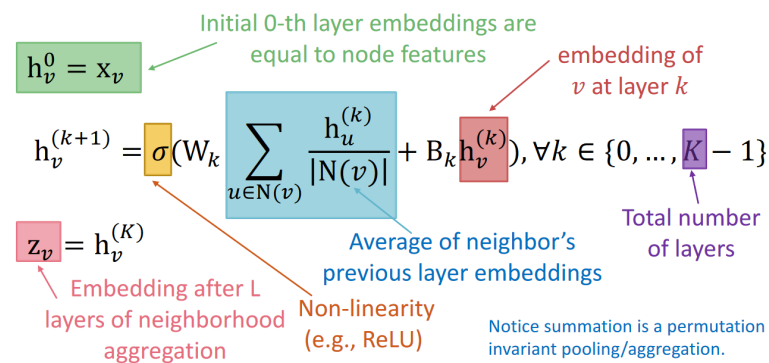
- The aggregation result is the same regardless of the ordering

- **Neighborhood aggregation:** Key distinctions are in how different approaches aggregate information across the layers



Transformation choice – average embedding

- Ordering invariant
- Examples
 - Average messages
 - Apply NN
 - Apply linear transformation
 - Follow by non-linearity
- Transform current layer features + aggregated previous child nodes messages



Model Parameters: W , B ; neighborhood aggregation, transformation

$$\begin{aligned} h_v^{(0)} &= x_v \\ h_v^{(k+1)} &= \sigma(W_k \sum_{u \in N(v)} \frac{h_u^{(k)}}{|N(v)|} + B_k h_v^{(k)}), \forall k \in \{0 \dots K-1\} \\ z_v &= h_v^{(K)} \end{aligned}$$

Trainable weight matrices (i.e., what we learn)

Final node embedding

We can feed these **embeddings into any loss function** and run SGD to **train the weight parameters**

- h_v^k : the hidden representation of node v at layer k
- W_k : weight matrix for neighborhood aggregation
- B_k : weight matrix for transforming hidden vector of self


Matrix Formulation - efficient sparse matrix

Averaging of neighbour embedding

- Many aggregations can be performed **efficiently** by (sparse) matrix operations
- Then
 - Node embedding is the average of the neighbour embedding
 - i.e. the adjacency matrix * embedding spaces at a layer
- in short
 - averaging of neighbour embedding
 - summing and averaging can be rewritten as
 - **matrix multiplication (dot product)**
 - the drawing is what correspond to what

- Let $H^{(k)} = [h_1^{(k)} \dots h_{|V|}^{(k)}]^T$
- Then: $\sum_{u \in N_v} h_u^{(k)} = A_{v,:} H^{(k)}$
- Let D be diagonal matrix where $D_{v,v} = \text{Deg}(v) = |N(v)|$
 - The inverse of D : D^{-1} is also diagonal: $D_{v,v}^{-1} = 1/|N(v)|$
- Therefore,

$$\sum_{u \in N(v)} \frac{h_u^{(k-1)}}{|N(v)|} \Rightarrow H^{(k+1)} = D^{-1} A H^{(k)}$$



Re-writing update function in matrix form

- In practice, this implies that **efficient sparse matrix** multiplication can be used (\tilde{A} is sparse)
- Note: not all GNNs can be expressed in matrix form, when aggregation function is **complex**

- Re-writing update function in matrix form:

$$H^{(k+1)} = \sigma(\tilde{A} H^{(k)} W_k^T + H^{(k)} B_k^T)$$

where $\tilde{A} = D^{-1} A$

$H^{(k)} = [h_1^{(k)} \dots h_{|V|}^{(k)}]^T$

- Red: neighborhood aggregation
- Blue: self transformation

How to train GNN: supervised & unsupervised setting

- Node embedding z_v is a function of input graph
- **Supervised setting:** we want to minimize the loss \mathcal{L} (see also Slide 15):

$$\min_{\Theta} \mathcal{L}(\mathbf{y}, f(\mathbf{z}_v))$$

- \mathbf{y} : node label
- \mathcal{L} could be L2 if \mathbf{y} is real number, or cross entropy if \mathbf{y} is categorical
- **Unsupervised setting:**
 - No node label available
 - Use the graph structure as the supervision!

Unsupervised Training

- No labels

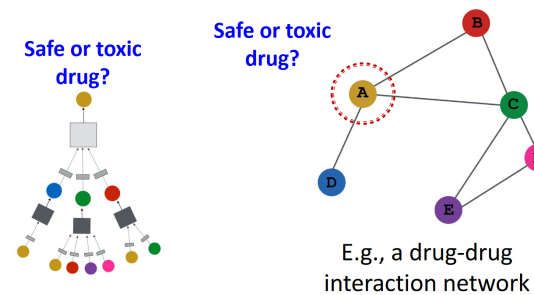
- “Similar” nodes have similar embeddings

$$\mathcal{L} = \sum_{z_u, z_v} \text{CE}(y_{u,v}, \text{DEC}(z_u, z_v))$$

- Where $y_{u,v} = 1$ when node u and v are **similar**
- CE** is the cross entropy (Slide 16)
- DEC** is the decoder such as inner product (Lecture 4)
- Node similarity** can be anything from Lecture 3, e.g., a loss based on:
 - Random walks** (node2vec, DeepWalk, struc2vec)
 - Matrix factorization**
 - Node proximity in the graph**

Supervised Training

- Directly train the model for a supervised task (e.g., node classification)




Cross entropy loss

- If label is 1, want output to be 1
- If 0, want 0


- Use cross entropy loss (Slide 16)

$$\mathcal{L} = \sum_{v \in V} y_v \log(\sigma(z_v^T \theta)) + (1 - y_v) \log(1 - \sigma(z_v^T \theta))$$




Model design

Step 1,2 - neighbour aggregation & loss function

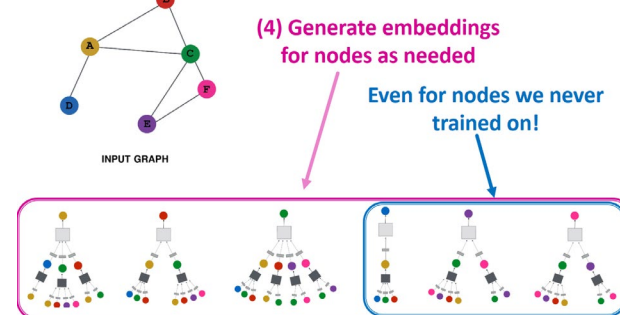


Step 3 – train on a batch



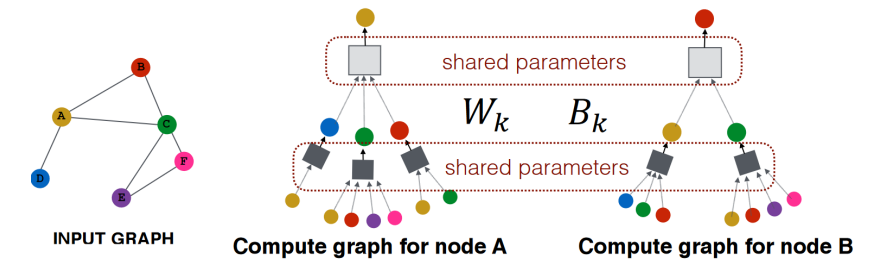
Step 4 - enerate embeddings for nodes as needed

- Generalisability.** Train embedding for one graph and transfer to another graph




Inductive capability

- The same aggregation parameters are **shared** for all nodes
- The number of model parameters is **sublinear** in $|V|$
- # model parameters **W & B**, depends on
 - #features / embedding dimensionality (since shared)
 - not the size of graph (#nodes)
- Thus, generalize to **unseen** nodes




New graph



Inductive node embedding → Generalize to entirely unseen graphs
E.g., train on protein interaction graph from model organism A and generate embeddings on newly collected data about organism B

New nodes

- One forward pass can generate new embedding for the new node as the graph evolving



- Many application settings constantly encounter previously unseen nodes:
 - E.g., Reddit, YouTube, Google Scholar
- Need to generate new embeddings “on the fly”

More

Permutation Invariance


- Graph does not have a canonical order of the nodes
- For a graph with m nodes, there are $m!$ different order plans.
- Are other neural network architectures, e.g., MLPs, permutation invariant / equivariant?
 - No
 - Switching the order of the input leads to different outputs!
 - This explains why the **naïve MLP** approach **fails** for **graphs**

- Consider we learn a function f that maps a graph $G = (A, X)$ to a vector \mathbb{R}^d then


$$f(A_1, X_1) = f(A_2, X_2)$$

A is the adjacency matrix
 X is the node feature matrix

Order plan 1: A_1, X_1



Order plan 2: A_2, X_2





For two order plans,
output of f should
be the same!

Equivariant Property


- Message passing and neighbor aggregation in graph convolution networks is **permutation equivariant**.
- The target node (blue) has the **same computation graph** for **different order plans**

equivariant.



GNNs subsume CNNs and Transformers

Convolutional neural network (CNN) layer with 3x3 filter:



$$\text{CNN formulation: } h_v^{(l+1)} = \sigma(\sum_{u \in N(v) \cup \{v\}} W_l^u h_u^{(l)}), \forall l \in \{0, \dots, L-1\}$$

$N(v)$ represents the 8 neighbor pixels of v .

GNN vs. CNN

- CNN can be seen as a special GNN with fixed neighbor size and ordering
 - The size of the filter is pre-defined for a CNN
 - The advantage of GNN is it processes arbitrary graphs with different degrees for each node
 - CNN is not permutation equivariant
 - Switching the order of pixels will leads to different outputs.

Convolutional neural network (CNN) layer with 3x3 filter:




- GNN formulation (previous slide): $h_v^{(l+1)} = \sigma(W_l \sum_{u \in N(v)} \frac{h_u^{(l)}}{|N(v)|} + B_l h_v^{(l)}), \forall l \in \{0, \dots, L-1\}$
- CNN formulation: $h_v^{(l+1)} = \sigma(\sum_{u \in N(v) \cup \{v\}} W_l^u h_u^{(l)}), \forall l \in \{0, \dots, L-1\}$
- if we rewrite: $h_v^{(l+1)} = \sigma(\sum_{u \in N(v)} W_l^u h_u^{(l)} + B_l h_v^{(l)}), \forall l \in \{0, \dots, L-1\}$

$$\text{GNN formulation: } h_v^{(l+1)} = \sigma(W_l \sum_{u \in N(v)} \frac{h_u^{(l)}}{|N(v)|} + B_l h_v^{(l)}), \forall l \in \{0, \dots, L-1\}$$

$$\text{CNN formulation: } h_v^{(l+1)} = \sigma(\sum_{u \in N(v)} W_l^u h_u^{(l)} + B_l h_v^{(l)}), \forall l \in \{0, \dots, L-1\}$$

Key difference: We can learn different W_l^u for different "neighbor" u for pixel v on the image. The reason is we can pick an order for the 9 neighbors using **relative position** to the center pixel: $\{(-1, -1), (-1, 0), (-1, 1), \dots, (1, 1)\}$

- Key component: **self-attention**
 - Every token/word attends to all the other tokens/words via matrix calculation
- Since each word attends to all the other words, the computation graph of a transformer layer is identical to that of a GNN on the fully-connected "word" graph.



Takeaway

- What if the ordering plan does matter?